



Heart Disease Prediction

using an Artificial Neural Network (ANN)

TEAM

23EG106C01	Aedula Sricharan
23EG106C02	Akkinepally Nagasai
23EG106C03	Akshith Sai Nalla

THE PROBLEM

What are we trying to solve?

Heart disease is a leading cause of death.

- Doctors look at many test results to judge risk.
- That takes time and expert skill.
- Early warning can save lives.

Our question:



**Can a computer
learn to spot
heart disease?**

In simple words

We give the computer past patient records and let it learn the pattern of 'healthy' vs 'has disease' — so it can judge a new patient.

THE DATA

What the model learns from

303

patient records
(UCI Cleveland dataset)

13 in → 1 out

13 health measurements
predict 1 answer: disease? yes / no

A few of the inputs:

Age

Sex

Chest pain type

Blood pressure

Cholesterol

Max heart rate

Blood sugar


ECG result

In simple words

Think of each patient as one row in an Excel sheet: 13 numbers describe them, and the last column says whether they had heart disease.

STEP 1 — PREPARING THE DATA

Cleaning before learning




In simple words

Raw data is messy. We tidy it so every feature is comparable, then keep a part of the data hidden to fairly test the model later.

STEP 2 — THE MODEL

A simple neural network




In simple words

Each layer mixes the numbers and passes them on, learning the hidden pattern. The last layer gives one score: higher = more likely to have heart disease.

STEP 3 — TRAINING

How the model learns



The idea

- Model makes a guess
- We measure the error
- It adjusts to do better
- Repeat many times

In simple words

The line going down means the model is making fewer mistakes each round. We stop automatically when it stops improving (no over-studying).

THE RESULTS

How well does it work?

82%

Accuracy


0.88

ROC-AUC


80%

Recall

(5-fold cross-validation —
a fair, stable score)



Confusion matrix



ROC curve

What it means

It catches 80% of real patients with disease.

0.88 AUC = it separates sick vs healthy well.

Confusion matrix shows correct vs wrong guesses.

CONCLUSION

What we built & learned

What we did

- Built a neural network in PyTorch
- Trained it on 303 patient records
- Tested it fairly with cross-validation
- Compared it to classic ML models

Key takeaways

- ~82% accuracy, 0.88 AUC
- Catches 80% of true cases
- On small data, simple models can match deep learning too

Next steps:

tune settings · explain predictions (SHAP) · try more data

Thank you!